

Contents

1	Introduction to Data Compression	2
1.1	Learning Objectives	2
1.2	Introduction and Motivation: Why Compress Data?	2
1.2.1	Benefits of Data Compression	3
1.3	Lossless vs. Lossy Compression	3
1.3.1	Lossless Compression	3
1.3.2	Lossy Compression	4
1.3.3	Choosing Between Lossless and Lossy	4
1.4	Compression Performance Metrics	4
1.4.1	Size-Based Metrics	4
1.4.2	Rate-Based Metrics	4
1.5	Worked Example: Audio Compression Metrics	5
1.6	Huffman Coding	5
1.6.1	Step-by-Step Huffman Coding Example	5
1.6.2	Key Properties of Huffman Coding	7
1.7	End of Chapter Questions	7
2	Theory of Compression — Limits and Optimality	11
2.1	Types of Codes: From Ambiguous to Instantaneous	11
2.2	Basic Terminology and Notation	12
2.3	Information and Redundancy: The Core Concepts	14
2.3.1	Information: A Formal Measure of Uncertainty Reduction	14
2.3.2	Redundancy: The Enemy of Information and the Friend of Compression	15
2.4	Entropy: The Fundamental Limit	16
2.4.1	What is Entropy? Different Perspectives	16
2.4.2	Calculating Entropy: Step by Step	16
2.4.3	Entropy of English Text: A Practical Case Study	17
2.4.4	Beyond First-Order Entropy: The Full Picture	17
2.4.5	Entropy Rate	18
2.4.6	The Entropy Theorem: Why It Matters	18
2.4.7	Key Takeaways	20
2.5	Entropy as a Lower Bound	21
2.5.1	The Fundamental Inequality	21
2.6	Kraft-McMillan Inequality: Core Theoretical Tool	21
2.6.1	Statement and Interpretation	21
2.7	Optimality of Huffman Codes (Theory Only)	22
2.7.1	The Optimality Theorem	22

2.7.2	Relation to Entropy	22
2.7.3	Why Huffman is Optimal but Not Perfect	23
2.8	Limitations of Symbol-by-Symbol Coding	23
2.8.1	Three Fundamental Limitations	23
2.9	Block Coding and Improved Efficiency	24
2.9.1	The Block Coding Idea	24
2.9.2	Key Mathematical Results	24
2.9.3	Step-by-Step Example	25
2.10	Trade-Offs in Block Coding	25
2.10.1	The Engineering Challenges	25
2.11	From Block Coding to Modern Compression	26
2.11.1	Arithmetic Coding: Fractional-Bit Block Coding	26
2.11.2	Context Modeling: Approximating Large Blocks	26
2.12	What Theory Guarantees vs What Practice Achieves	26
2.13	Summary and Key Takeaways	26
2.13.1	The Big Picture	27
2.13.2	Looking Forward	27
3	Advanced Entropy Coding & Extensions	28
3.1	Introduction & Motivation	28
3.2	Coding Taxonomy & Framework	28
3.3	Shannon Coding (1948)	30
3.4	Shannon-Fano Coding (1949)	31
3.5	Canonical Huffman Codes	32
3.6	Adaptive Huffman Coding	33
3.7	Arithmetic Coding: The Paradigm Shift	34
3.8	Comparison & Synthesis	35
3.9	Forward Look	36
4	Source Modeling and Statistical Dependence	40
4.1	Introduction: Beyond Coding	40
4.2	Memoryless vs. Sources with Memory	41
4.3	Conditional Entropy and Mutual Information	42
4.4	Markov Sources	43
4.5	Entropy Rate Revisited	44
4.6	Context Modeling in Practice	44
4.7	Case Study: Text Compression Modeling	45
4.8	The Modeling–Coding Separation Principle	46
4.9	Adaptive vs. Static Modeling	47
4.10	Summary and Forward Look	48

Data Compression

Lecture Notes

Dr. Faisal Aslam

1: Lecture 1: Introduction to Data Compression

1.1 Learning Objectives

By the end of this lecture, students will be able to:

- Understand the motivation and benefits of data compression
- Differentiate between lossless and lossy compression techniques
- Compute and interpret common compression performance metrics
- Apply the Huffman coding algorithm step by step
- Analyze real-world compression trade-offs

1.2 Introduction and Motivation: Why Compress Data?

Data compression is the process of representing information using fewer bits than its original form. It is a fundamental component of modern computing systems, enabling efficient storage, faster communication, and reduced operational costs.

Everyday applications of compression include:

- Streaming audio and video
- Image storage and sharing
- File archiving and backups
- Network communication and cloud services

Definition

Data Compression is the process of reducing the number of bits required to represent information, either:

- **Losslessly**: allowing exact reconstruction of the original data
- **Lossily**: allowing controlled loss of information to achieve higher compression

1.2.1 Benefits of Data Compression

Data compression provides three key benefits that are critical in modern computing:

1. Reduce Storage Space:

- Allows more data to be stored in the same physical space
- Enables archival of historical data that would otherwise be discarded
- Reduces hardware requirements for storage systems

2. Reduce Communication Time and Bandwidth:

- Enables faster file transfers and downloads
- Makes high-quality streaming (4K/8K video) practical over limited bandwidth
- Reduces latency in real-time applications like video conferencing and online gaming
- Allows IoT devices to transmit data efficiently over wireless networks

3. Save Money:

- Reduces cloud hosting costs (storage and egress fees)
- Lowers communication costs for data transmission
- Decreases capital expenditure on storage hardware
- Reduces energy consumption for data centers and network infrastructure

1.3 Lossless vs. Lossy Compression

1.3.1 Lossless Compression

Lossless compression guarantees perfect reconstruction of the original data. It is essential when accuracy and data integrity are critical.

Typical applications:

- Text files and source code
- Executables and databases
- Medical, scientific, and legal data

1.3.2 Lossy Compression

Lossy compression achieves higher compression ratios by discarding information that is less perceptible or less important.

Typical applications:

- Audio (MP3, AAC)
- Images (JPEG)
- Video (H.264, H.265)

1.3.3 Choosing Between Lossless and Lossy

Factor	Lossless Compression	Lossy Compression
Reconstruction	Exact	Approximate
Data sensitivity	High	Moderate to low
Typical ratios	Low to moderate	High
Quality impact	None	Controlled degradation

Table 1: Lossless vs. Lossy Compression

1.4 Compression Performance Metrics

1.4.1 Size-Based Metrics

$$\text{Compression Ratio (CR)} = \frac{\text{Original Size}}{\text{Compressed Size}}$$
$$\text{Compression Factor} = \frac{\text{Compressed Size}}{\text{Original Size}}$$
$$\text{Space Savings (\%)} = \left(1 - \frac{\text{Compressed Size}}{\text{Original Size}}\right) \times 100\%$$

Interpretation:

- Larger compression ratios indicate better compression
- Smaller compression factors indicate better compression

1.4.2 Rate-Based Metrics

$$\text{Bits per Sample (bps)} = \frac{\text{Compressed Size (bits)}}{\text{Number of samples}}$$
$$\text{Bit-rate (bps)} = \frac{\text{Compressed Size (bits)}}{\text{Time (seconds)}}$$

These metrics are particularly important in audio and video compression systems.

1.5 Worked Example: Audio Compression Metrics

Example

Uncompressed Audio Properties

- Duration: 180 seconds
- Sampling rate: 44.1 kHz
- Bit depth: 16 bits
- Channels: 2 (stereo)

Original Size Calculation

$$\text{Total samples} = 180 \times 44,100 \times 2 = 15,876,000$$

$$\text{Size (bits)} = 15,876,000 \times 16 = 254,016,000$$

$$\text{Size (MB)} = \frac{254,016,000}{8 \times 1,048,576} \approx 30.27$$

Compression Results

Method	Size (MB)	CR	Savings	Bit-rate
FLAC (lossless)	18.16	1.67:1	40%	807 kbps
MP3 @ 320 kbps	6.75	4.49:1	77.7%	320 kbps
AAC @ 256 kbps	5.40	5.61:1	82.2%	256 kbps

1.6 Huffman Coding

Huffman coding is a widely used **lossless compression algorithm** that assigns variable-length binary codes to symbols based on their frequencies. More frequent symbols receive shorter codes.

1.6.1 Step-by-Step Huffman Coding Example

Example

Message: MISSISSIPPI RIVER (17 characters including space)

Symbol Frequencies

Symbol	Frequency
I	5
S	4
P	2
R	2
M	1
V	1
E	1
(space)	1

Tree Construction

1. Combine M(1) + V(1) → 2
2. Combine E(1) + (space)(1) → 2
3. Combine P(2) + R(2) → 4
4. Combine 2 + 2 → 4
5. Combine 4 + 4 → 8
6. Combine I(5) + S(4) → 9
7. Combine 8 + 9 → 17

One Possible Code Assignment

Symbol	Code	Length
I	00	2
S	01	2
P	100	3
R	101	3
M	1100	4
V	1101	4
E	1110	4
(space)	1111	4

Compressed Size

$$5(2) + 4(2) + 2(3) + 2(3) + 4(1) = 52 \text{ bits}$$

Original Size (ASCII) = $17 \times 8 = 136$ bits

Compression Ratio = $136/52 \approx 2.62 : 1$

1.6.2 Key Properties of Huffman Coding

- Produces prefix-free codes
- Enables instantaneous decoding
- Guarantees minimum average code length among prefix codes
- Widely used in practical compression systems

1.7 End of Chapter Questions

Exercise 1.0

Problem 1: Basic Compression Metrics

An uncompressed grayscale image has the following properties:

- Resolution: 1024×1024 pixels
- Bit depth: 8 bits per pixel

After compression, the image size is 320 KB.

Calculate:

- Original image size in KB
- Compression ratio
- Compression factor
- Space savings percentage

Exercise 1.1

Problem 2: Audio Bit-rate and Storage

A mono audio recording has the following parameters:

- Duration: 5 minutes
- Sampling rate: 48 kHz
- Bit depth: 16 bits

The file is compressed using a lossy codec to a constant bit-rate of 192 kbps.

Calculate:

- Size of the uncompressed audio file in MB
- Size of the compressed file in MB

- (c) Compression ratio
- (d) Bits per sample after compression

Exercise 1.2

Problem 3: Comparing Compression Options

A video clip has an uncompressed data rate of 120 Mbps. Three compression options are available:

Option	Compressed Bit-rate
A	6 Mbps
B	3 Mbps
C	1.5 Mbps

For each option, calculate:

- (a) Compression ratio
- (b) Data consumed for a 10-minute video (in MB)

Which option would you choose for:

- (i) Live video streaming?
- (ii) Archival storage?

Briefly justify your answers.

Exercise 1.3

Problem 4: Huffman Coding Construction

Given the following symbol frequencies:

Symbol	Frequency
A	10
B	8
C	6
D	5
E	4
F	3
G	2
H	2

- (a) Construct the Huffman tree step by step
- (b) Assign a binary code to each symbol

- (c) Compute the total number of bits required to encode the message
- (d) Calculate the average number of bits per symbol

Exercise 1.4

Problem 5: Fixed-Length vs. Huffman Coding

Using the symbol set from Problem 4:

- (a) Determine the minimum fixed-length code required
- (b) Compute the total number of bits using fixed-length coding
- (c) Compare the result with Huffman coding
- (d) Calculate the percentage reduction in total bits achieved by Huffman coding

Exercise 1.5

Problem 6: Text Compression Scenario

A text file contains 50,000 characters and is stored using 8-bit ASCII encoding.

After compression using a lossless algorithm, the file size becomes 18 KB.

Calculate:

- (a) Original file size in KB
- (b) Compression ratio
- (c) Compression factor
- (d) Space savings percentage

Explain why compression ratios for text files vary significantly depending on content.

Exercise 1.6

Problem 7: Practical Design Question

You are designing a compression system for a wearable health-monitoring device that:

- Records sensor data continuously
- Has limited storage capacity
- Requires exact data reconstruction
- Operates on a low-power processor

- (a) Should the system use lossless or lossy compression? Explain.
- (b) Which performance metrics are most important in this scenario?
- (c) Would a variable-length coding scheme be appropriate? Why or why not?

2: Lecture 2: Theory of Compression — Limits and Optimality

2.1 Types of Codes: From Ambiguous to Instantaneous

Definition

Types of Codes

- **Non-singular Code:** Each source symbol maps to a distinct codeword

$$x_i \neq x_j \Rightarrow C(x_i) \neq C(x_j)$$

- **Uniquely Decodable Code:** Every finite sequence of codewords corresponds to exactly one sequence of source symbols

$$C(x_1)C(x_2)\cdots C(x_n) = C(y_1)C(y_2)\cdots C(y_m) \Rightarrow n = m \text{ and } x_i = y_i$$

- **Prefix Code (Instantaneous Code):** No codeword is a prefix of another codeword

$$\forall i \neq j : \quad C(x_i) \text{ is not a prefix of } C(x_j)$$

Key Relationships:

$$\text{Prefix Codes} \subset \text{Uniquely Decodable Codes} \subset \text{Non-singular Codes}$$

Important

Why Prefix Codes are Special

- **Instantaneous decoding:** Can decode as soon as codeword ends (no lookahead needed)
- **Tree representation:** Always correspond to leaves of a binary tree
- **Kraft inequality:** Always satisfy $\sum 2^{-\ell_i} \leq 1$
- **Practical:** Used in Huffman coding, many real-world compressors

Example

Example: Comparing Different Code Types

For symbols $\{A, B, C, D\}$ with probabilities $\{0.5, 0.25, 0.125, 0.125\}$:

Code Type	A	B	C	D	Property
Non-singular	0	1	00	11	Distinct but ambiguous: "00" = AA or C?
Uniquely decodable	0	01	011	0111	Unique but need lookahead
Prefix code	0	10	110	111	Instant decoding: "0" = A, stop
Optimal prefix	0	10	110	111	Also Huffman optimal

Decoding examples:

- **Prefix code "010110"**: 0→A, 10→B, 110→C = "ABC" (instant)
- **Uniquely decodable "00111"**: Need to scan ahead to determine split
- **Non-singular "00"**: Ambiguous! Could be "AA" or "C"

Key insight: Prefix codes sacrifice some flexibility in codeword lengths (must satisfy Kraft inequality) for the benefit of instantaneous decoding.

2.2 Basic Terminology and Notation

Definition

Alphabet

An *alphabet* \mathcal{X} is a finite set of possible symbols. Examples:

- Binary alphabet: $\mathcal{X} = \{0, 1\}$
- English letters: $\mathcal{X} = \{A, \dots, Z\}$
- Bytes: $\mathcal{X} = \{0, 1, \dots, 255\}$

Definition

Symbol

A *symbol* is a single element drawn from an alphabet. For example, the letter E is a symbol from the English alphabet.

Definition

Random Variable

A *random variable* X is a function that assigns a symbol or value to each outcome in a sample space:

$$X : \Omega \rightarrow \mathcal{X}$$

- Ω : Sample space (e.g., all possible states of a data source)

- \mathcal{X} : Set of possible values (alphabet, e.g., $\{0, 1\}$, ASCII characters)
- For each $\omega \in \Omega$, $X(\omega)$ is the value assigned to outcome ω

Example: Binary Source

- $\Omega = \{\text{emits 0}, \text{emits 1}\}$ (or could be more complex underlying physics)
- $\mathcal{X} = \{0, 1\}$
- $X(\text{emits 0}) = 0, X(\text{emits 1}) = 1$
- Probabilities: $P(X = 0) = P(\{\omega : X(\omega) = 0\}) = p, P(X = 1) = 1 - p$

Why this matters for compression:

- The entropy $H(X)$ depends on the probability distribution induced by X
- For $x \in \mathcal{X}$: $P(X = x) = P(\{\omega \in \Omega : X(\omega) = x\})$
- $H(X) = -\sum_{x \in \mathcal{X}} P(X = x) \log_2 P(X = x)$

Definition

Source

A *source* is a process that generates a sequence of symbols (X_1, X_2, X_3, \dots) according to some probability law. In this lecture, we assume discrete sources unless stated otherwise.

Definition

Message (or Sequence)

A *message* is a finite sequence of symbols generated by the source:

$$x^n = (x_1, x_2, \dots, x_n)$$

Compression algorithms operate on messages, not on individual symbols.

Definition

Code and Codewords

A *code* assigns a binary string (codeword) to each symbol or message.

- Source symbols \rightarrow codewords (e.g., Huffman coding)
- Messages \rightarrow bitstreams (e.g., arithmetic coding)

Definition

Block Length

The *block length* n is the number of source symbols grouped together and encoded as a unit. Larger block lengths generally allow better compression but increase delay and complexity.

Definition

Model

A *model* estimates the probabilities of symbols or sequences. Better models lead to better compression by reducing uncertainty.

2.3 Information and Redundancy: The Core Concepts

2.3.1 Information: A Formal Measure of Uncertainty Reduction

In information theory, information is defined rigorously as a **quantitative measure of the reduction in uncertainty** that results from observing the outcome of a random event.

Definition 2.1. Let X be a random event that occurs with probability $p = \Pr(X)$. The **information content** (or self-information) $I(X)$ provided by the occurrence of X is defined as:

$$I(X) = \log_b \left(\frac{1}{p} \right) = -\log_b(p)$$

where:

- $b = 2$ yields **bits** (binary digits)
- $b = e$ yields **nats** (natural units)
- $b = 10$ yields **hartleys** or **dits**

Example

Predictability vs. Information:

- In a city where it rains every day, the statement “It rained today” conveys almost no information because it was expected
- A file that contains only the bit ‘1’ provides very little information
- A coin that always lands heads produces outcomes, but no information

Key idea: Perfect predictability implies zero information gain.

Example

Daily Weather Forecast — Information Content:

- Sunny in Phoenix (probability 0.9): $I = -\log_2 0.9 \approx 0.15$ bits
- Snow in Phoenix (probability 0.001): $I = -\log_2 0.001 \approx 9.97$ bits
- Rain in Seattle (probability 0.3): $I = -\log_2 0.3 \approx 1.74$ bits

Interpretation: Rare events carry more information because they reduce uncertainty the most.

2.3.2 Redundancy: The Enemy of Information and the Friend of Compression

Redundancy refers to predictable or repeated structure in data. It is what allows data to be represented using fewer bits.

1. **Spatial Redundancy:** Neighboring data values are highly correlated

Example

In a photograph of a clear blue sky, most neighboring pixels have nearly identical color values.

- **Naive:** Store the RGB value of each pixel independently
- **Smarter:** Encode repeated pixel values using run-length encoding
- **Even smarter:** Predict each pixel from its neighbors and encode only the small prediction error

2. **Statistical Redundancy:** Some symbols occur far more frequently than others

Example

English letter frequencies:

Letter	Frequency	Letter	Frequency
E	12.7%	Z	0.07%
T	9.1%	Q	0.10%
A	8.2%	J	0.15%

Frequent letters get shorter codes in variable-length coding schemes.

3. **Knowledge Redundancy:** Information already known to both encoder and decoder

4. **Perceptual Redundancy:** Information that humans cannot perceive

2.4 Entropy: The Fundamental Limit

2.4.1 What is Entropy? Different Perspectives

Definition

Shannon Entropy of a discrete random variable X with possible values $\{x_1, x_2, \dots, x_n\}$ having probabilities $\{p_1, p_2, \dots, p_n\}$:

$$H(X) = - \sum_{i=1}^n p_i \log_2 p_i \text{ bits}$$

Two Complementary Interpretations:

1. **Average Information Content:** Expected value of information content across all symbols
2. **Uncertainty or Surprise:** Measures how uncertain we are about the next symbol

2.4.2 Calculating Entropy: Step by Step

Example

Binary Source Example - Detailed Calculation:

Consider a biased coin: $P(\text{Heads}) = 0.8$, $P(\text{Tails}) = 0.2$

Step 1: Calculate individual information content:

$$I_H = -\log_2(0.8) \approx 0.3219 \text{ bits}$$

$$I_T = -\log_2(0.2) \approx 2.3219 \text{ bits}$$

Step 2: Calculate entropy as expected value:

$$H = 0.8 \times 0.3219 + 0.2 \times 2.3219 = 0.7219 \text{ bits}$$

Step 3: Verify using direct formula:

$$H = -[0.8 \log_2(0.8) + 0.2 \log_2(0.2)] \approx 0.7219 \text{ bits}$$

Key Insights:

- **Extreme cases:**

- Fair coin ($P=0.5$): $H = 1.0$ bit (maximum uncertainty)
- Always heads ($P=1.0$): $H = 0$ bits (no uncertainty)
- 90% heads: $H \approx 0.469$ bits

2.4.3 Entropy of English Text: A Practical Case Study

Example

Calculating English Letter Entropy:

Based on letter frequencies in typical English text:

$$H \approx 4.18 \text{ bits/letter}$$

Layered Interpretation:

- First-order entropy (letters independent): 4.18 bits/letter
- Actual uncertainty is lower: Letters have dependencies ($Q \rightarrow U$)
- Comparison with encoding schemes:

Encoding Method	Bits/Letter
Naive (5 bits for 26 letters)	5.00
Huffman (letter-based)	4.30
Using digram frequencies	3.90
Using word frequencies	2.30
Optimal with full context	~ 1.50

2.4.4 Beyond First-Order Entropy: The Full Picture

Higher-Order Entropies quantify uncertainty while accounting for increasing context:

- Zero-order entropy (H_0): $H_0 = \log_2 |\mathcal{X}|$
- First-order entropy (H_1): $H_1 = -\sum p(x) \log_2 p(x)$
- Second-order entropy (H_2): $H_2 = -\sum p(x, y) \log_2 p(x|y)$
- Nth-order entropy (H_N): $H_N = -\sum p(x_1, \dots, x_N) \log_2 p(x_N | x_1, \dots, x_{N-1})$

2.4.5 Entropy Rate

The **entropy rate** of a source is defined as the limiting uncertainty per symbol when arbitrarily long contexts are available:

$$H_\infty = \lim_{N \rightarrow \infty} H_N$$

2.4.6 The Entropy Theorem: Why It Matters

Definition

Expected Code Length

For a source with symbols $\{x_1, x_2, \dots, x_n\}$ having probabilities $\{p_1, p_2, \dots, p_n\}$, and a code that assigns codeword lengths $\{\ell_1, \ell_2, \dots, \ell_n\}$, the **expected code length** L is:

$$L = \mathbb{E}[\ell(X)] = \sum_{i=1}^n p_i \ell_i \quad (\text{bits per symbol})$$

This measures the average number of bits needed to encode one symbol from the source.

Definition

Compression Ratio and Efficiency

For a source with entropy $H(X)$ and code with expected length L :

- **Compression ratio:** $\rho = \frac{\text{original bits}}{\text{compressed bits}}$
- **Efficiency:** $\eta = \frac{H(X)}{L} \leq 1$
- **Redundancy:** $R = L - H(X) \geq 0$

Perfect compression occurs when $\eta = 1$ (100% efficient) and $R = 0$.

Important

Shannon's Source Coding Theorem (1948)

For a discrete memoryless source with entropy H and any $\epsilon > 0$:

1. **Converse (Impossibility Result):**

No lossless coding scheme can achieve expected code length $L < H$.

$L \geq H$ for any uniquely decodable code
--

2. **Achievability (Possibility Result):**

There exists a lossless coding scheme (specifically, block coding with suffi-

ciently large block size n) such that:

$$H \leq L < H + \epsilon$$

Equivalently: For any $\epsilon > 0$, $\exists n$ such that:

$$\frac{L_n}{n} < H + \epsilon$$

where L_n is the expected length for blocks of size n .

Interpretation:

- **Entropy is the fundamental limit:** H bits/symbol is the best we can ever do
- **We can get arbitrarily close:** With clever coding, we can approach this limit as closely as desired
- **The gap is achievable:** The " $+\epsilon$ " represents practical overhead that can be made arbitrarily small

Example

Understanding the Theorem with Numbers

Consider a binary source with $p(0) = 0.9$, $p(1) = 0.1$:

$$H = -0.9 \log_2 0.9 - 0.1 \log_2 0.1 \approx 0.469 \text{ bits/symbol}$$

- **Naive coding:** Use 1 bit per symbol $\rightarrow L = 1.0$, efficiency $\eta = 0.469/1.0 = 46.9\%$
- **Huffman coding:** $0 \rightarrow 0$, $1 \rightarrow 1$ (same as naive!) $\rightarrow L = 1.0$, $\eta = 46.9\%$ *Why so bad?* Because we're coding symbols individually.
- **Block coding (n=2):** Code pairs of symbols:

$$\begin{aligned} 00 &\rightarrow 0 \quad (\ell = 1, p = 0.81) \\ 01 &\rightarrow 10 \quad (\ell = 2, p = 0.09) \\ 10 &\rightarrow 110 \quad (\ell = 3, p = 0.09) \\ 11 &\rightarrow 111 \quad (\ell = 3, p = 0.01) \end{aligned}$$

$L_2 = 0.81 \times 1 + 0.09 \times 2 + 0.09 \times 3 + 0.01 \times 3 = 1.29$ bits/block Per symbol:
 $L = L_2/2 = 0.645$ bits/symbol, $\eta = 0.469/0.645 \approx 72.7\%$

- **Block coding ($n=3$):** Would get even closer to 0.469
- **Theoretical limit:** As $n \rightarrow \infty, L \rightarrow 0.469$

Key insight: The theorem tells us:

1. We can never beat 0.469 bits/symbol (impossibility)
2. We can get as close as we want to 0.469 bits/symbol (achievability)

Important

What the Theorem Does NOT Say

- **It doesn't say how to construct the code** - just that one exists
- **It doesn't guarantee practical implementation** - block size n might need to be huge
- **It doesn't account for computational complexity** - the code might be too complex to implement
- **It assumes we know the true probabilities** - in practice, we estimate them

Yet, this theorem is revolutionary because it:

1. Establishes a **fundamental limit** (like the speed of light in physics)
2. Provides a **benchmark** for evaluating compression algorithms
3. Guides algorithm design toward this limit

2.4.7 Key Takeaways

- Entropy measures both **average information** and **uncertainty**
- Higher-order models reduce entropy by exploiting dependencies
- The entropy rate represents the ultimate compression limit
- Shannon's theorem precisely separates the *possible* from the *impossible*

2.5 Entropy as a Lower Bound

2.5.1 The Fundamental Inequality

Theorem 2.2 (Entropy Lower Bound). *For any **uniquely decodable** code C for source X :*

$$L(C) \geq H(X)$$

where $L(C) = \mathbb{E}[\ell(X)] = \sum_i p_i \ell_i$ is the expected code length.

Example

Binary Source with $p(0) = 0.9, p(1) = 0.1$

$$H = -0.9 \log_2 0.9 - 0.1 \log_2 0.1 \approx 0.469 \text{ bits/symbol}$$

Why we can't achieve 0.4 bits/symbol:

1. For 100 symbols, typical sequences: $2^{100 \times 0.469} \approx 2^{46.9}$
2. To encode all uniquely, need at least $2^{46.9}$ codewords
3. At 0.4 bits/symbol, total bits = 40
4. # codewords $\leq 2^{40} < 2^{46.9} \rightarrow$ impossible!

2.6 Kraft-McMillan Inequality: Core Theoretical Tool

2.6.1 Statement and Interpretation

Theorem 2.3 (Kraft-McMillan Inequality (Binary Case)). *Let $\ell_1, \ell_2, \dots, \ell_m$ be the lengths of codewords in a **prefix code**. Then:*

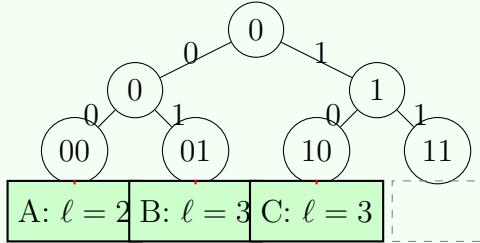
$$\sum_{i=1}^m 2^{-\ell_i} \leq 1$$

Conversely, if integers ℓ_1, \dots, ℓ_m satisfy this inequality, then there exists a binary prefix code with these lengths.

Example

Tree Visualization of Kraft Inequality

Consider a binary tree of depth $L = \max \ell_i$:



Calculating Kraft sum for $\{\ell_A = 2, \ell_B = 3, \ell_C = 3\}$:

$$\sum 2^{-\ell_i} = 2^{-2} + 2^{-3} + 2^{-3} = 0.25 + 0.125 + 0.125 = 0.5 \leq 1$$

Example

Testing Code Feasibility

1. Valid lengths: $\{1, 2, 3, 3\}$

$$\sum = 2^{-1} + 2^{-2} + 2^{-3} + 2^{-3} = 0.5 + 0.25 + 0.125 + 0.125 = 1.0 \quad \text{VALID}$$

2. Invalid lengths: $\{1, 1, 2\}$

$$\sum = 2^{-1} + 2^{-1} + 2^{-2} = 0.5 + 0.5 + 0.25 = 1.25 > 1 \quad \text{INVALID}$$

2.7 Optimality of Huffman Codes (Theory Only)

2.7.1 The Optimality Theorem

Theorem 2.4 (Huffman Optimality). *Given a source with symbol probabilities p_1, p_2, \dots, p_m , the Huffman algorithm produces a prefix code that **minimizes** the expected code length $L = \sum_{i=1}^m p_i \ell_i$ among all prefix codes.*

2.7.2 Relation to Entropy

For any Huffman code:

$$H(X) \leq L_{\text{Huffman}} < H(X) + 1$$

Example

Understanding the "+1" Gap

Consider source with probabilities $\{0.6, 0.3, 0.1\}$:

$$H \approx 1.295 \text{ bits}$$

Ideal (non-integer) lengths: $-\log_2 p_i = \{0.737, 1.737, 3.322\}$

Huffman code: 0.6→0, 0.3→10, 0.1→11

$$L = 0.6 \times 1 + 0.3 \times 2 + 0.1 \times 2 = 1.4 \text{ bits}$$

Comparison:

- Entropy: 1.295 bits
- Huffman: 1.400 bits
- Gap: 0.105 bits (much less than 1!)

2.7.3 Why Huffman is Optimal but Not Perfect

Important

Huffman is Optimal Within a Restricted Class

Huffman is optimal among:

- **Symbol-by-symbol** codes
- **Prefix** codes
- **Static** codes

But real optimality might require:

- **Block coding**
- **Fractional bits** (arithmetic coding)
- **Adaptive probabilities**

2.8 Limitations of Symbol-by-Symbol Coding

2.8.1 Three Fundamental Limitations

1. **Cannot Exploit Dependencies**
2. **Integer Length Constraint:** $\ell_i \in \mathbb{Z}^+$ but $-\log_2 p_i \in \mathbb{R}$
3. **Memoryless Assumption**

Example

English Text: The Cost of Symbol-by-Symbol

- **First-order entropy** (ignoring dependencies): 4.0 bits/letter
- **Actual entropy rate** (with dependencies): 1.5 bits/letter
- **Huffman on letters**: 4.0 bits/letter
- **Gap**: 2.5 bits/letter wasted due to ignoring dependencies

2.9 Block Coding and Improved Efficiency

2.9.1 The Block Coding Idea

Instead of coding symbols individually, group them into blocks of length n :

$$\mathbf{X} = (X_1, X_2, \dots, X_n)$$

Definition

n th Extension of a Source

For a source with alphabet \mathcal{X} , the n th extension has alphabet:

$$\mathcal{X}^n = \{(x_1, \dots, x_n) : x_i \in \mathcal{X}\}$$

with size $|\mathcal{X}|^n$.

2.9.2 Key Mathematical Results

Theorem 2.5 (Entropy of Block Source). *For a discrete memoryless source:*

$$H(X^n) = nH(X)$$

Theorem 2.6 (Block Coding Performance). *There exists a prefix code C_n for X^n such that:*

$$nH(X) \leq L_n < nH(X) + 1$$

Dividing by n :

$$H(X) \leq \frac{L_n}{n} < H(X) + \frac{1}{n}$$

Important

The Magic of Block Coding

As $n \rightarrow \infty$:

$$\frac{L_n}{n} \rightarrow H(X)$$

We can approach entropy **arbitrarily closely** by making blocks larger!

2.9.3 Step-by-Step Example

Example

Binary Source: $p(0) = 0.9, p(1) = 0.1, H \approx 0.469$

Step 1: $n = 1$ (**symbol-by-symbol**)

- Huffman: $0 \rightarrow 0, 1 \rightarrow 1$
- $L_1 = 1$ bit/symbol
- Efficiency: $\eta = 0.469/1 = 46.9\%$

Step 2: $n = 2$ (**code pairs**)

- Block probabilities: $P(00) = 0.81, P(01) = 0.09, P(10) = 0.09, P(11) = 0.01$
- Codes: $00 \rightarrow 0, 01 \rightarrow 10, 10 \rightarrow 110, 11 \rightarrow 111$
- Per symbol: $L_2/2 = 0.645$ bits/symbol
- Efficiency: $\eta = 0.469/0.645 = 72.7\%$

2.10 Trade-Offs in Block Coding

2.10.1 The Engineering Challenges

1. **Exponential Alphabet Growth:** $|\mathcal{X}^n| = |\mathcal{X}|^n$
2. **Memory Requirements:** Huffman tree has $2m^n - 1$ nodes
3. **Computational Complexity:** $O(m^n \log m^n)$
4. **Delay and Latency:** Must wait for n symbols

2.11 From Block Coding to Modern Compression

2.11.1 Arithmetic Coding: Fractional-Bit Block Coding

Important

Arithmetic Coding as "Infinite Block Coding"

Arithmetic coding cleverly avoids the exponential growth problem:

- **Idea:** Encode entire message as a single real number in $[0,1)$
- **No explicit blocks:** Processes symbols sequentially
- **Fractional bits:** Achieves $L \approx H(X)$ without large n
- **Removes integer constraint:** No "+1" overhead!

2.11.2 Context Modeling: Approximating Large Blocks

Instead of explicit block coding, modern compressors use:

1. **Context Models:** Predict next symbol based on previous k symbols
2. **Prediction + Residual Coding:** Encode only prediction error
3. **Dictionary Methods (LZ family):** Build dictionary of previously seen phrases

2.12 What Theory Guarantees vs What Practice Achieves

Aspect	Theory Guarantees	Practice Achieves
Optimality	Can approach entropy arbitrarily closely	Gets close, but with practical limits
Block Size	$n \rightarrow \infty$ gives optimality	n limited by memory, latency, complexity
Complexity	Ignored, infinite resources allowed	Critical constraint; often dominates design

2.13 Summary and Key Takeaways

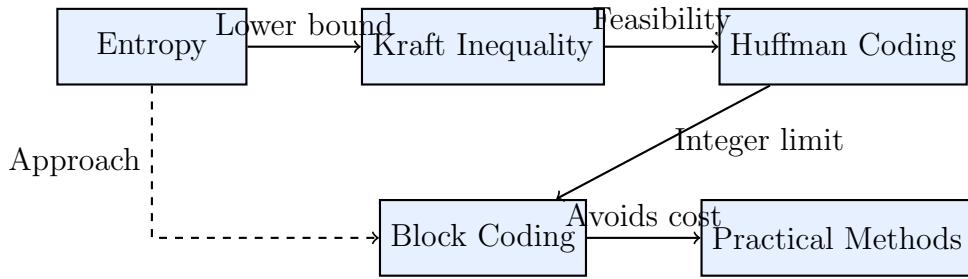
Important

Five Fundamental Lessons

1. **Entropy is the Absolute Limit:** $L \geq H(X)$ for any lossless code

2. Kraft-McMillan Constrains All Codes: $\sum 2^{-\ell_i} \leq 1$
3. Huffman is Optimal Among Prefix Codes: But limited by integer lengths
4. Block Coding Allows Approaching Entropy: $\lim_{n \rightarrow \infty} \frac{L_n}{n} = H(X)$
5. Practical Compression Balances Efficiency and Complexity

2.13.1 The Big Picture



2.13.2 Looking Forward

- Next lecture: Arithmetic Coding: Removes integer constraint
- Then: Dictionary Methods (LZ family): Adaptive to data statistics
- Finally: Modern Compressors: Combining multiple techniques

Final Thought

Shannon's 1948 paper told us *exactly how good compression could possibly be*. Every compressor since has been trying to approach that limit while staying within practical constraints.

The gap between theory and practice is where engineering creativity lives!

3: Lecture 3: Advanced Entropy Coding & Extensions

Lecture 3: Beyond Huffman – Advanced Entropy Coding Methods

3.1 Introduction & Motivation

Important

Recall Huffman Coding Limitations:

- **Integer code lengths:** Cannot reach entropy bound for highly skewed distributions
- **Static vs. Adaptive:** Standard Huffman requires prior knowledge of probabilities
- **Codebook overhead:** Need to transmit/store the coding tree
- **Symbol-by-symbol constraint:** Processes one symbol at a time

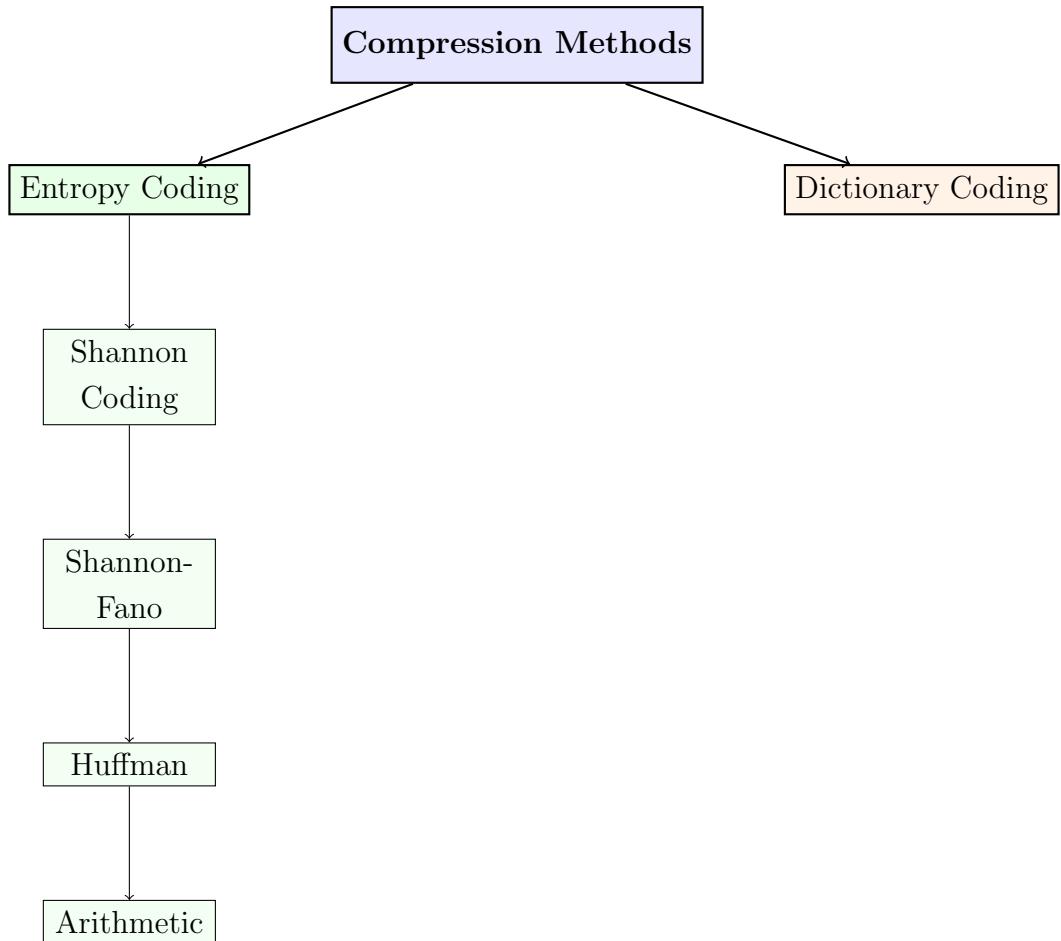
Lecture Roadmap:

1. **Framework:** Coding taxonomy and conceptual organization
2. **Historical methods:** Shannon & Shannon-Fano coding
3. **Practical improvements:** Canonical and Adaptive Huffman
4. **Next generation:** Arithmetic coding paradigm
5. **Synthesis:** Comparison and forward look

3.2 Coding Taxonomy & Framework

Definition

Coding Taxonomy: Classification of compression methods based on key characteristics



Key Dimensions in Coding Taxonomy:

1. **Modeling vs. Coding Separation:**
 - **Modeling:** Probability estimation of symbols
 - **Coding:** Assigning bit sequences based on probabilities
2. **Static vs. Adaptive vs. Universal:**
 - **Static:** Fixed probabilities known in advance
 - **Adaptive:** Probabilities updated during encoding
 - **Universal:** Adapts to unknown distributions
3. **Block Codes vs. Stream Codes:**
 - **Block:** Fixed-length input blocks
 - **Stream:** Continuous data processing
4. **Symbol-by-Symbol vs. Message-Wide:**
 - **Symbol-by-symbol:** Each symbol coded independently (Huffman)
 - **Message-wide:** Entire message coded as one unit (Arithmetic)

3.3 Shannon Coding (1948)

Definition

Shannon Coding: A constructive method derived from Shannon's source coding theorem:

$$l_i = \lceil -\log_2 p_i \rceil$$

where p_i is the probability of symbol i .

Example

Example: Given symbols with probabilities:

Symbol	Probability	$-\log_2 p_i$
A	0.5	1.0
B	0.25	2.0
C	0.125	3.0
D	0.125	3.0

Shannon code construction:

1. Calculate lengths: $l_A = \lceil 1.0 \rceil = 1$, $l_B = 2$, $l_C = 3$, $l_D = 3$
2. Sort by probability (decreasing)
3. Assign cumulative probabilities: $F_A = 0$, $F_B = 0.5$, $F_C = 0.75$, $F_D = 0.875$
4. Convert F_i to binary with l_i bits:
 - A: $0.0_2 \rightarrow 0$
 - B: $0.5_{10} = 0.1_2 \rightarrow 10$
 - C: $0.75_{10} = 0.11_2 \rightarrow 110$
 - D: $0.875_{10} = 0.111_2 \rightarrow 111$

Expected length: $L = 0.5 \times 1 + 0.25 \times 2 + 0.125 \times 3 + 0.125 \times 3 = 1.75$ bits/symbol

Important

Properties of Shannon Coding:

- **Constructive proof:** Demonstrates Kraft-McMillan inequality
- **Simple to compute:** Direct from probabilities
- **Not optimal:** Unlike Huffman, doesn't minimize expected length
- **Theoretical importance:** Foundation for Shannon's theorem

3.4 Shannon-Fano Coding (1949)

Definition

Shannon-Fano Coding: A top-down recursive splitting method developed independently by Shannon and Fano.

Algorithm 1 Shannon-Fano Coding Algorithm

Require: Symbols with probabilities $p_1 \geq p_2 \geq \dots \geq p_n$

Ensure: Prefix code for each symbol

```
1: Sort symbols by decreasing probability
2: function SF(symbols)
3: if  $|symbols| = 1$  then
4:   return (assign empty code)
5: end if
6: Split symbols into two subsets  $S_1$  and  $S_2$  with nearly equal total probabilities
7: Append '0' to codes of symbols in  $S_1$ 
8: Append '1' to codes of symbols in  $S_2$ 
9: SF( $S_1$ )
10: SF( $S_2$ )
11: end function
```

Example

Example: Same symbols as before:

Symbol	Probability
A	0.5
B	0.25
C	0.125
D	0.125

Construction:

1. Split 1: $\{A\}(0.5)$ vs $\{B, C, D\}(0.5) \rightarrow A:0$, others:1
2. Split 2: $\{B\}(0.25)$ vs $\{C, D\}(0.25) \rightarrow B:10$, others:11
3. Split 3: $\{C\}(0.125)$ vs $\{D\}(0.125) \rightarrow C:110$, D:111

Codes: A=0, B=10, C=110, D=111

Expected length: Same as Huffman in this case: 1.75 bits/symbol

Important

Historical Note: Shannon-Fano coding was developed before Huffman coding (1952). While simpler conceptually, it is **not always optimal**. Huffman's bottom-

up approach guarantees optimality for given probabilities.

3.5 Canonical Huffman Codes

Definition

Canonical Huffman Code: A Huffman code where codes are assigned in a specific canonical (standard) form, enabling more efficient representation and faster decoding.

Algorithm 2 Canonical Huffman Construction

Require: Symbol code lengths l_i from standard Huffman

Ensure: Canonical Huffman codes

- 1: Sort symbols by: (1) code length l_i , (2) symbol value
 - 2: Assign first code: $code = 0$ (in binary with l_{min} bits)
 - 3: **for** each symbol in sorted order **do**
 - 4: Assign current $code$ to symbol
 - 5: Increment $code$ by 1
 - 6: **if** next symbol has longer code length **then**
 - 7: Left-shift $code$ by difference in lengths
 - 8: **end if**
 - 9: **end for**
-

Example

Example: Suppose Huffman gives these lengths:

Symbol	Length	Original Huffman
A	2	00
B	3	010
C	3	011
D	3	100
E	4	1010
F	4	1011

Canonical construction:

1. Sort: A(2), B(3), C(3), D(3), E(4), F(4)
2. Start: A gets code 00 (binary 0 in 2 bits)
3. Increment: B gets 010 (binary 2 in 3 bits: 010)
4. Increment: C gets 011 (binary 3: 011)
5. Increment: D gets 100 (binary 4: 100)

6. For E: length increases to 4, so left-shift: $100 \rightarrow 1000$
7. Increment: F gets 1001

Canonical codes: A=00, B=010, C=011, D=100, E=1000, F=1001

Important

Advantages of Canonical Huffman:

- **Compact representation:** Only need to store code lengths, not the full tree
- **Fast decoding:** Use table-based lookup instead of tree traversal
- **Standardized:** Used in DEFLATE (ZIP, gzip), JPEG, PNG, MPEG

Transmission: Send only: \langle symbol count, lengths \rangle instead of full tree

3.6 Adaptive Huffman Coding

Definition

Adaptive Huffman Coding: Dynamically updates the Huffman tree as symbols are processed, requiring only one pass over the data.

Key Algorithms:

- **FGK Algorithm** (Faller, Gallager, Knuth, 1978)
- **Vitter's Algorithm** (1987) - more efficient

Important

How it works:

1. Start with empty tree or initial estimates
2. For each symbol:
 - Encode using current tree
 - Update symbol frequency
 - Reconstruct/update tree incrementally
3. Sibling property maintained for optimality

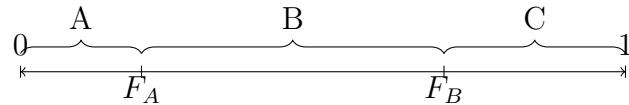
Applications: Early UNIX compress utility, network protocols where distribution changes.

3.7 Arithmetic Coding: The Paradigm Shift

Definition

Arithmetic Coding: Encodes an entire message into a single fractional number in the interval $[0, 1)$, approaching the entropy bound very closely.

Core Idea: Represent messages as subintervals of $[0, 1)$:



Algorithm 3 Arithmetic Encoding Algorithm

Require: Message $m = s_1 s_2 \dots s_k$, symbol probabilities p_i

Ensure: Final interval $[low, high)$

```

1:  $low \leftarrow 0.0$ ,  $high \leftarrow 1.0$ 
2: for each symbol  $s$  in  $m$  do
3:    $range \leftarrow high - low$ 
4:    $high \leftarrow low + range \times F_s$  { $F_s$ : cumulative prob up to  $s$ }
5:    $low \leftarrow low + range \times F_{s-1}$  { $F_{s-1}$ : cumulative prob before  $s$ }
6: end for
7: return any number in  $[low, high)$ 

```

Example

Example: Encode message "CAB" with probabilities: A(0.5), B(0.25), C(0.25)

Symbol	Probability	Cumulative
A	0.5	0.5
B	0.25	0.75
C	0.25	1.0

Encoding:

1. Start: $[0, 1)$
2. Process 'C': $[0.75, 1.0)$ (C occupies $[0.75, 1.0)$)
3. Process 'A': $range = 0.25$
 - $low = 0.75 + 0.25 \times 0.0 = 0.75$
 - $high = 0.75 + 0.25 \times 0.5 = 0.875$
 - New interval: $[0.75, 0.875)$
4. Process 'B': $range = 0.125$
 - $low = 0.75 + 0.125 \times 0.5 = 0.8125$

- $high = 0.75 + 0.125 \times 0.75 = 0.84375$
- Final interval: $[0.8125, 0.84375)$

Output: Any number in $[0.8125, 0.84375)$, e.g., 0.8125 in binary

Important

Practical Implementation Issues:

- **Finite precision:** Use integer arithmetic with scaling
- **Renormalization:** Output bits when interval confined to one half
- **Carry-over:** Handle when interval spans midpoint
- **Termination:** Need special end-of-message symbol

Adaptive Arithmetic Coding: Easier than adaptive Huffman - just update probabilities as you go!

3.8 Comparison & Synthesis

Method	Optimal?	Adaptive?	Complexity	Near Entropy?	Key Applications
Shannon Coding	No	No	Low	No	Theoretical proofs
Shannon-Fano	No	No	Low	Sometimes	Historical
Huffman	Yes*	No	Low	Moderate	General purpose
Canonical Huffman	Yes*	No	Low	Moderate	DEFLATE, JPEG, PNG
Adaptive Huffman	Yes*	Yes	Medium	Moderate	Early compressors
Arithmetic Coding	Near-opt	Yes	High	Yes (close)	JPEG2000, H.264, HEVC

*Optimal for symbol-by-symbol coding given probabilities

Important

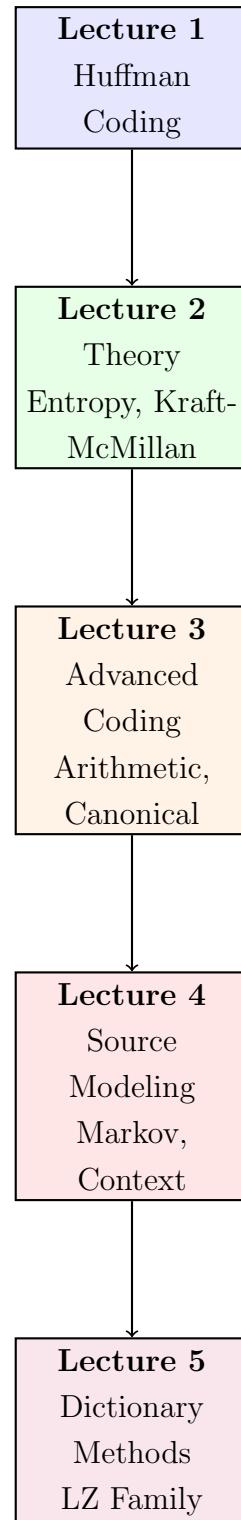
Key Insights:

- **Huffman vs. Arithmetic:** Huffman is simpler but has an "integer penalty"; Arithmetic approaches entropy bound

- **Modern standards:** Arithmetic coding (CABAC) used in video compression for 10-20% better compression
- **Practical choice:** For general compression, Canonical Huffman (DEFLATE); for media, Arithmetic coding
- **The missing piece:** All these methods assume we have good probability estimates. Where do those come from?

3.9 Forward Look

The Complete Picture: What Comes Next?



Next Lecture: Source Modeling and Statistical Dependence

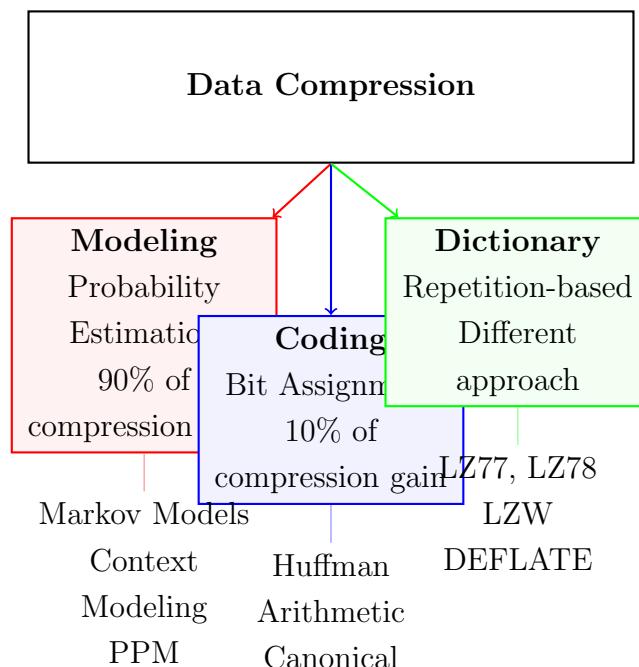
- **The missing half:** We now know how to code efficiently, but where do the probabilities come from?
- **Real data has memory:** 'Q' is usually followed by 'U' in English text
- **Markov models:** Capturing dependencies between symbols

- **Context modeling:** Using past symbols to predict future ones
- **The modeling-coding separation:** Modern compressors separate these two tasks

The Big Question for Next Time:

If Arithmetic coding can get within 0.01 bits of entropy,
what's the real limit to compression?
The answer: It's not the coding, it's the *modeling*!

The Two Pillars of Compression (Revised View):



Exercise 3.0

Exercise 3.1: Given symbols with probabilities: A(0.4), B(0.3), C(0.2), D(0.1)

- Construct a Shannon code and compute its expected length
- Construct a Shannon-Fano code
- Compare with Huffman code from Lecture 2

Exercise 3.1

Exercise 3.2: Convert the following Huffman code to canonical form:

Symbol	Huffman Code
A	0
B	100
C	101
D	110
E	1110
F	1111

Exercise 3.2

Exercise 3.3: Encode the message "ABAC" using arithmetic coding with probabilities: A(0.6), B(0.3), C(0.1). Show each step.

Exercise 3.3

Exercise 3.4: Thinking Ahead: Consider the English phrase "THE QUICK BROWN FOX"

- (a) If we use Huffman coding with letter frequencies, what's wrong with this approach?
- (b) How might knowing that 'Q' is usually followed by 'U' help compression?
- (c) Why would arithmetic coding be better than Huffman for this kind of data?

End of Lecture 3 – Advanced Entropy Coding Methods

Next: Lecture 4 – Source Modeling and Statistical Dependence

We now have efficient coding methods. Next: Where do the probabilities come from?

4: Lecture 4: Source Modeling and Statistical Dependence

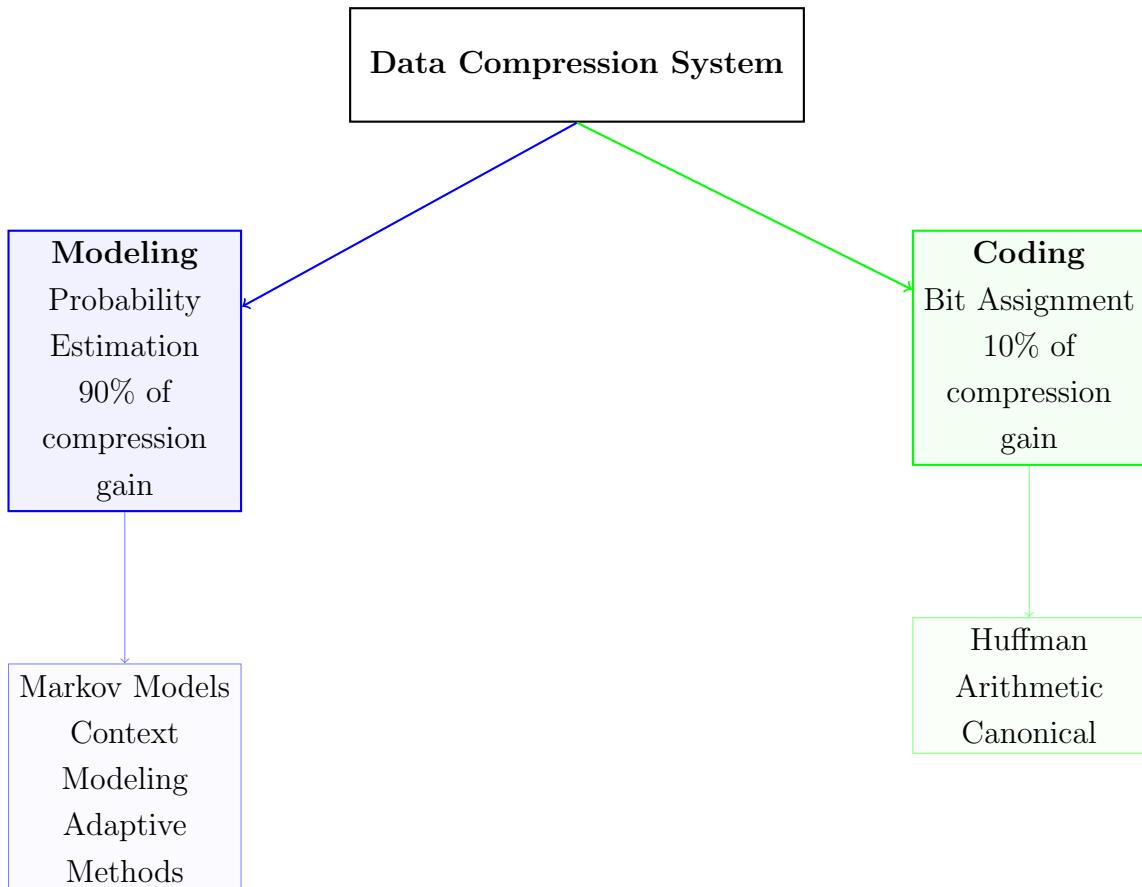
Lecture 4: Beyond Coding – The Power of Source Modeling

4.1 Introduction: Beyond Coding

Important

Key Insight: In the first three lectures, we've focused on **coding** - efficient ways to represent symbols given their probabilities. Now we address the other half: **modeling** - how to get good probability estimates in the first place.

The Big Picture:



Why Real Data Defies IID Assumptions:

- **IID (Independent Identically Distributed):** Assumption behind simple Huffman
- **Reality:** Data has **memory** and **dependencies**
- Example: In English text, 'Q' is almost always followed by 'U'

- Example: In images, neighboring pixels are highly correlated

Today's Roadmap:

1. Understand statistical dependence in data
2. Learn Markov models for capturing memory
3. Explore context modeling techniques
4. See practical examples with real data
5. Connect modeling to coding (Lecture 3)

4.2 Memoryless vs. Sources with Memory

Definition

Memoryless Source (IID): Each symbol is generated independently of all previous symbols. Probability distribution: $P(X_n = x) = p(x)$ for all n .

Definition

Source with Memory: The probability of a symbol depends on previous symbols. Example: $P(X_n = x | X_{n-1} = y, X_{n-2} = z, \dots)$.

Example

Examples of Dependence in Real Data:

- **Text:** 'TH' is common, 'TQ' is rare
- **Images:** Neighboring pixels have similar colors
- **Audio:** Sound waves have temporal continuity
- **Video:** Consecutive frames are nearly identical
- **Source code:** Keywords, variable names repeat

Important

Measuring Dependence: Autocorrelation

$$\rho(k) = \frac{\mathbb{E}[(X_t - \mu)(X_{t+k} - \mu)]}{\sigma^2}$$

where k is the lag. High $\rho(k)$ means strong dependence at distance k .

4.3 Conditional Entropy and Mutual Information

Definition

Conditional Entropy: The average uncertainty about X given knowledge of Y :

$$H(X|Y) = - \sum_{x \in \mathcal{X}} \sum_{y \in \mathcal{Y}} p(x, y) \log_2 p(x|y)$$

Example

Intuition: "Knowing the Past Helps Predict the Future"

Consider English letters:

- Unconditional: $H(\text{letter}) \approx 4.07$ bits
- Given previous letter: $H(\text{letter}|\text{previous}) \approx 3.36$ bits
- Given previous 2 letters: $H(\text{letter}|\text{previous 2}) \approx 2.77$ bits

Each additional letter of context reduces uncertainty!

Definition

Mutual Information: Measures how much knowing Y reduces uncertainty about X :

$$I(X; Y) = H(X) - H(X|Y) = H(Y) - H(Y|X)$$

Example

Worked Example: English Letter Dependence

Let X = current letter, Y = previous letter.

$$\begin{aligned} H(X) &= 4.07 \text{ bits} \\ H(X|Y) &= 3.36 \text{ bits} \\ I(X; Y) &= 4.07 - 3.36 = 0.71 \text{ bits} \end{aligned}$$

This means knowing the previous letter gives us 0.71 bits of information about the current letter.

4.4 Markov Sources

Definition

Markov Property (Memory- m): The future depends only on the last m symbols:

$$P(X_n = x | X_{n-1}, X_{n-2}, \dots, X_1) = P(X_n = x | X_{n-1}, \dots, X_{n-m})$$

First-Order Markov Model ($m = 1$):

- Only the immediately previous symbol matters
- Represented by transition probabilities $p_{ij} = P(X_n = j | X_{n-1} = i)$
- Can be shown as a state diagram or transition matrix

Example

Example: Weather Prediction Markov Chain

States: {Sunny (S), Rainy (R)}

Transition probabilities:

	S	R
S	0.8	0.2
R	0.3	0.7

Interpretation: If today is sunny, 80% chance tomorrow is sunny, 20% chance rainy.

Higher-Order Markov Models:

- Order- k : Depends on last k symbols
- More accurate but exponentially more parameters
- Number of parameters grows as $|\mathcal{A}|^{k+1}$ where \mathcal{A} is alphabet size

Important

Memory-Complexity Trade-off:

- **Order 0:** 26 parameters for English (simple but weak)
- **Order 1:** $26 \times 26 = 676$ parameters
- **Order 2:** $26 \times 26 \times 26 = 17,576$ parameters
- **Order 5:** $26^6 \approx 308$ million parameters!

This is the **context explosion problem**.

4.5 Entropy Rate Revisited

Definition

Entropy Rate of a Stationary Source:

$$H(\mathcal{X}) = \lim_{n \rightarrow \infty} \frac{1}{n} H(X_1, X_2, \dots, X_n)$$

For a stationary Markov chain of order m :

$$H(\mathcal{X}) = H(X_{m+1} | X_1, \dots, X_m)$$

Example

Calculating Entropy Rate for First-Order Markov Chain:

For weather example with stationary distribution $\pi = [\pi_S, \pi_R]$:

$$H(\mathcal{X}) = \pi_S H(X|S) + \pi_R H(X|R)$$

where:

$$H(X|S) = -0.8 \log_2 0.8 - 0.2 \log_2 0.2 \approx 0.7219$$

$$H(X|R) = -0.3 \log_2 0.3 - 0.7 \log_2 0.7 \approx 0.8813$$

$$\pi_S = 0.6, \quad \pi_R = 0.4 \quad (\text{solve } \pi P = \pi)$$

$$H(\mathcal{X}) = 0.6 \times 0.7219 + 0.4 \times 0.8813 \approx 0.788 \text{ bits}$$

Important

What This Means for Compression:

- **IID assumption:** Limit = $H(X)$ (e.g., 4.07 bits/letter for English)
- **With modeling:** Limit = $H(\mathcal{X})$ (e.g., 2.3 bits/letter for English)
- **Potential gain:** Up to 45% better compression!

4.6 Context Modeling in Practice

Definition

Context Modeling: Maintain separate probability distributions for each possible context (history).

Fixed-Length vs. Variable-Length Contexts:

- **Fixed-length:** Always use last k symbols as context

- **Variable-length:** Use longest matching context in database
- Example: PPM (Prediction by Partial Matching) uses variable-length

Example

The Context Explosion Problem:

For English text (26 letters + space):

Order	Contexts	Parameters
0	1	27
1	27	729
2	729	19,683
3	19,683	531,441
4	531,441	14,348,907
5	14,348,907	387,420,489

By order 5: 387 million parameters need estimation!

Solutions to Context Explosion:

1. **Escaping:** Fall back to lower-order model when context unseen
2. **Blending:** Combine predictions from different order models
3. **Pruning:** Remove low-frequency contexts
4. **Adaptive methods:** Update probabilities as data arrives

4.7 Case Study: Text Compression Modeling

Example

English Text Compression with Different Models:

Model Type	Bits/Letter	Compression vs. ASCII
ASCII (baseline)	8.00	0%
Order-0 (Huffman)	4.07	49%
Order-1 (Bigram)	3.36	58%
Order-2 (Trigram)	2.77	65%
Order-3	2.43	70%
Order-5 (PPM)	2.23	72%
Optimal (Shannon 1951)	1.3	84%

Note: Each step improves compression by better modeling!

PPM (Prediction by Partial Matching):

- Uses **variable-length** contexts

- Tries highest-order model first
- Escapes to lower order if context unseen
- Blends probabilities from different orders
- State-of-the-art for text compression in 1990s

Important

Practical Entropy Reduction:

IID model (Huffman) : 4.07 bits/letter

With context modeling : 2.23 bits/letter

Improvement : 45% better compression!

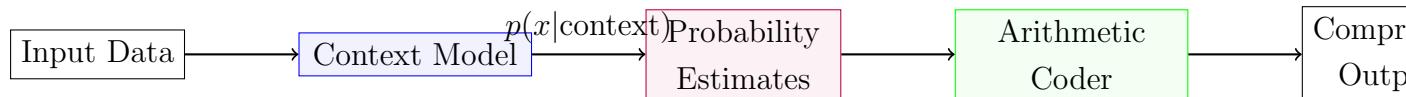
4.8 The Modeling–Coding Separation Principle

Definition

Modeling–Coding Separation: Modern compressors separate probability estimation (modeling) from bit assignment (coding).

Historical Evolution:

- **Early:** Integrated (Huffman builds tree from frequencies)
- **Modern:** Separated (Model → Probabilities → Arithmetic Coder)



Example

How PPM + Arithmetic Beats Huffman:

1. **PPM:** Sees context "TH" → predicts E with 80% probability
2. **Arithmetic:** Encodes E using $p = 0.8 \rightarrow 0.32$ bits
3. **Huffman:** Would need at least 1 bit for any symbol
4. **Gain:** 0.32 bits vs 1+ bits = 3× better for this symbol!

Important

Why Arithmetic Coding is the Perfect Backend:

- Can handle **fractional bits** per symbol
- Accepts **changing probabilities** symbol by symbol
- Works with **adaptive models** naturally
- Achieves entropy bound for good models

4.9 Adaptive vs. Static Modeling

Definition

Static Models: Train once on representative data, use fixed model for all files.

- **Pros:** Fast encoding/decoding
- **Cons:** Model may not match specific file
- **Example:** Early text compressors using English statistics

Definition

Semi-Adaptive (Two-Pass): First pass: collect statistics; Second pass: encode.

- **Pros:** Tailored to specific file
- **Cons:** Need to transmit model (overhead)
- **Example:** Standard Huffman with tree transmission

Definition

Fully Adaptive (One-Pass): Update model while encoding/decoding.

- **Pros:** No model transmission, adapts to local changes
- **Cons:** Slower, initial poor compression
- **Example:** Adaptive Huffman, PPM with update

Example

Comparison in Practice:

Type	Compression	Speed	Memory
Static	Medium	Fast	Low
Semi-Adaptive	Good	Medium	Medium
Fully Adaptive	Best	Slow	High

Choice depends on application constraints!

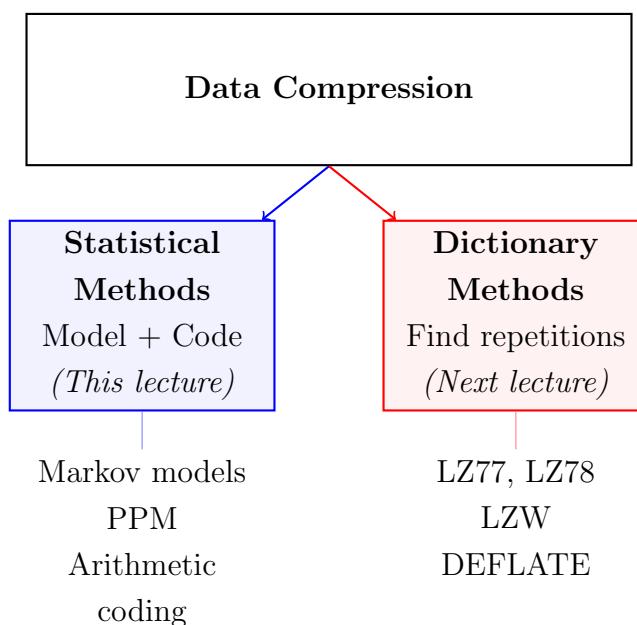
4.10 Summary and Forward Look

Important

Key Takeaways:

1. **The real compression is in modeling**, not just coding
2. **Context matters**: Using past symbols reduces uncertainty
3. **Markov models** capture memory in data
4. **Context explosion** limits practical model order
5. **Arithmetic coding** enables efficient use of good models
6. **Adaptive methods** avoid model transmission overhead

The Two Pillars Revisited:



Preview: Next Lecture on LZ Family (Dictionary Methods):

- A completely different approach: find repeating patterns

- No probability estimation needed
- Works well for files with exact repetitions
- Used in ZIP, GIF, PDF, and many modern formats
- Often combined with statistical methods in practice

Exercise 4.0

Exercise 4.1: Given the first-order Markov chain for weather:

	S	R
S	0.7	0.3
R	0.4	0.6

- Find the stationary distribution $\pi = [\pi_S, \pi_R]$
- Calculate the entropy rate $H(\mathcal{X})$
- How does this compare to an IID source with $P(S) = 0.55, P(R) = 0.45$?

Exercise 4.1

Exercise 4.2: Consider the text fragment: "THE CAT SAT ON THE MAT" a) Build an order-1 (bigram) model for this text b) Calculate $H(X)$ (order-0 entropy) c) Calculate $H(X|Y)$ where Y is previous letter (order-1 conditional entropy) d) How much mutual information exists between consecutive letters?

Exercise 4.2

Exercise 4.3: Explain why arithmetic coding is better than Huffman coding when used with: a) A high-order Markov model b) An adaptive context model c) A model that gives very skewed probabilities (e.g., $p = 0.99$ for one symbol)

End of Lecture 4 – Source Modeling and Statistical Dependence

Next: Dictionary-based Compression (LZ Family)